



Quality engineering for the AI era

AI × quality engineering · Industry research



A shared quality capability




Strategic vision · outcomes are ambitions to validate, not reported bank results

The constraints are organizational

What constrains adoption

Share of respondents reporting each barrier · World Quality Report 2025–26



Overlapping responses, not shares of a whole. Survey: >2,000 executives; question-specific n unavailable in the public release. Source: Capgemini, Sogeti and OpenText, November 2025. [\[W01\]](#)

Task gains depend on the setting

CUI ET AL. / THREE FIELD EXPERIMENTS

+26.08%

completed tasks

Developer sample

4,867

Precision

SE 10.3 percentage points

Setting

Coding assistants in three firms

METR / EARLY-2025 TRIAL

+19%

completion time

Developer / task sample

16 / 246

Precision

95% CI: +2% to +39%

Setting

Experienced developers; familiar repos


The unit, task and user population determine what transfers.

Separate outcomes; do not pool these percentages or treat them as QE savings rates.

Capacity and realized value

\$45k illustrative net economic impact
3.30% capacity released · 0.45% of addressable spend

Annual value bridge
\$ thousands · per \$10M addressable QA spend



Illustrative assumptions, not a forecast. Capacity is effort equivalent; cash savings need Finance-approved capture.

Adoption 50% · net task saving 20% · capacity captured 50%.

Activity share 55% · eligibility 60% · AI/pilot cost 1.0% · quality allowance 0.2%.

The industry direction and its limits

SOURCE / EVIDENCE TYPE	WHAT IT SAYS	STRATEGIC IMPLICATION
Gartner · forecast	90% of enterprise software engineers using AI code assistants by 2028.	Plan for widespread assisted work; tool use alone does not establish quality or savings.
McKinsey · survey	Nearly 300 leaders surveyed; 100 assessed impact. Top performers reported broader delivery improvements.	Invest in workflow and organizational change. These self-reported outcomes do not isolate causality.
DORA · observational research	AI interacts with the surrounding engineering system.	Improve feedback, testing and platform foundations alongside adoption.

Public sources, reviewed September 2026. Gartner’s gated testing reports contribute public abstracts only; no proprietary vendor ranking is reproduced. [G01][G02][G03][M01][D01]

Industrial cases reveal reusable patterns

INDUSTRIAL EXAMPLE	OBSERVED APPROACH	CAPABILITY TO DEVELOP
Meta · TestGen-LLM / ACH	Generated tests face build, reliability and coverage checks; ACH adds fault-oriented validation.	An independent test-quality gate, with meaningful assertions.
Google · Auto-Diagnose	Failure summaries link likely causes to relevant log lines in the review workflow.	Diagnosis with inspectable evidence and reviewer feedback.
Uber / UT Austin · FlakyGuard	Targeted execution context supports flaky-test repair in an enterprise Go repository.	Repair validation that preserves the original test intent.

Company-specific cases demonstrate mechanisms. They do not establish a transferable bank savings rate or a product comparison. [E04][E05][E06][E07]

The first workflows need a clear quality check

Easy to verify

BOUNDED STARTING POINTS

Test drafts with an executable contract
Failure summaries linked to logs

CONTROLLED EXECUTION

Sandbox test execution
Draft pull requests through scoped tools

Needs expert judgment

BUILD THE EVALUATION CAPABILITY

Risk analysis and exploratory test ideas
Expert review defines useful output

SEPARATE AUTHORITY DECISION

Production changes and customer actions
Require action-level evidence and recovery

Advice / draft artifacts


Greater action authority

What the findings change


RESEARCH OBSERVATION

DESIGN RESPONSE

EVIDENCE TO RETAIN




Quality improves through feedback



Proposed application-level assurance · evaluate again when context, behavior or authority changes

Ownership across the quality lifecycle




The quality role gains new disciplines

CAPABILITY	WHAT PRACTITIONERS NEED TO DO	OBSERVABLE LEARNING EVIDENCE
Test intent and oracles	Define expected behavior and challenge plausible but weak assertions.	A test catches a relevant fault and preserves the contract.
Evaluation and curation	Build representative cases; calibrate expert labels and model judges.	Evaluation results survive a held-out scenario review.
Agent assurance	Inspect tool authority, context boundaries and recovery behavior.	A denied-action exercise leaves no unauthorized side effect.
Workflow ownership	Measure review and correction work; coach teams through real tasks.	Task-level outcomes include the cost of verification.

Proposed capability agenda informed by research on verification work and lifecycle assurance; it is not a staffing reduction plan. [D02][G04][A01][E05]

The platform outlasts a model



Shared foundations and domain expertise

PRODUCT-SPECIFIC ASSETS

Payments

Transaction invariants
Failure and recovery cases

Customer service AI

Grounded response criteria
Escalation and tool scenarios

Engineering workflows

Repository test contracts
Review and rework baselines



REUSABLE ASSURANCE SERVICES


Dataset registry · Evaluation runners · Identity + gateway · Evidence store · Recovery tooling



REPLACEABLE MODEL AND TOOL ADAPTERS


Procure capabilities against local tests; preserve access to configurations, artifacts and evidence.

Evidence gates the next horizon




Proposed horizons, not a dated forecast or measured maturity score

The leadership evidence contract



Every readout includes task mix, sample, period, costs and uncertainty. No organizational results are claimed here.

Governance follows the system lifecycle



OSFI context: July 2026 bulletin = sound practices; revised E-23 takes effect 1 May 2027.

The strategic choices ahead

Ambition

Quality as a shared capability

Choose the delivery constraints and AI risks the organization needs to address.

Investment

Foundations and people

Develop evaluation skills, platform integration and accountable ownership.

Expansion

Value with evidence

Grow the capabilities that improve outcomes within their control boundary.

Vision: AI contributes more work while the organization strengthens its ability to judge and govern that work.

Choose the shared assurance model

01

Local assistants

Fast individual adoption
Duplicated context and checks
Local learning stays fragmented

Use to learn tasks

02

Shared assurance

Portable tests and evidence
Domain-owned expectations
Common evaluation and controls

Recommended strategic center

03

Bounded agents

Multi-step delegated work
Stronger action authority
Higher recovery obligations


Earn through evidence

Sequencing choice: standardize the proof before expanding delegated authority.


Trade-off: shared services require ownership; agent breadth requires independently tested containment.

What changes in a payment-API workflow

CURRENT WORKING HYPOTHESIS · VALIDATE WITH THE BANK



PROPOSED TARGET · NEGATIVE PAYMENT AMOUNTS MUST BE REJECTED



Shared: identity, evaluation runners, evidence schema. Domain: payment rules, test cases, acceptance and service outcomes.

Sequence shared assets before wider agency

HORIZON	SHARED INSTITUTIONAL ASSET	DOMAIN RESPONSIBILITY / TRADE-OFF
Prove the workflow	Task identity, approved models, baseline measurement and evidence contract.	Payment owner defines behavior; delivery team measures review effort. Keep scope narrow.
Reuse the assurance	Evaluation runners, policy enforcement and reusable failure cases.	Own domain cases, acceptance and outcomes. Shared services need funded owners.
Delegate bounded actions	Scoped tool adapters, recovery exercises and operational evidence.	Service owner accepts consequence and recovery duty. Expand one authority boundary at a time.

Client adoption requires funded foundations

Adoption decision

Required evidence accepted · owner accountable · foundation work funded



Business & people

Expected outcomes · reviewer capacity · offshore handoffs · measured baseline

Runnable engineering

Infrastructure · CI/CD · test data · dependency control · testable applications

Supported services

Frameworks & devices · AI access · run evidence · platform ownership

Authored adoption assumptions. Validate each application and workflow; unknown requirements remain open. [Assumptions and evidence](#) →

Fund the missing capabilities and reviewer capacity before promising scale. A workflow can progress only when its required evidence is accepted.

QE modernization and AI adoption

01

Trusted work

Reviewed scenarios, expected outcomes and ownership

AI drafts with domain review

A reviewer can identify a wrong answer

02

Repeatable execution

Reliable suites, isolated data and controlled dependencies

AI drafts tests that run in CI

Another engineer reproduces the result

03

Shared QE services

Reusable environments, contracts, artifacts and support

Assistance works across teams

A second team onboards and operates it

04

Broader AI execution

Evaluated tools, bounded actions and recovery

Agents execute within approved scope

Failures stop, evidence persists, people can intervene

Proposed capability progression. Reviewed assistance can begin while foundations improve. Broader execution requires evidence for its specific scope. [Modernization workstreams and sources](#)

Develop reliable, reusable QE capability alongside reviewed AI assistance. Broader execution depends on the evidence available for the selected workflow.

Financial-services outcomes

Libra Internet Bank

35% less time

Test-creation time

Index: previous time = 100



UiPath customer case. Normalized from the reported reduction; review effort is unspecified. [F1](#)

Separate outcomes and denominators. Each panel has its own meaning; no combined savings rate is calculated.


Reported customer outcomes use different measures and evidence types. Use them to select a trial; establish the client’s own baseline.

Goldman Sachs

36% → 72%

Unit-test coverage

Share of a selected module (%)




Diffblue customer case. Coverage measures exercised code; it does not establish financial correctness. [F2](#)

Fiserv

65% fewer incidents

Major incidents

Index: previous year = 100



Tricentis modernization case. AI testing was a later pilot; this is not an AI-attributed reduction. [F4](#)

A focused first engagement

Requirements ready

Onboarding design

Draft scenarios in the current test-management workflow.

Proof

Accepted coverage and measured review effort.

Proposed entry points. Choose one using the client's constraints; [inspect prerequisites and prepare a trial brief](#).

Select the entry point that fits the client’s foundations. Build reusable assets and retain evidence for the next adoption decision.

Execution repeatable

Payment API tests

Generate candidates in the existing framework and CI pipeline.

Proof

Correct behavior passes; a known fault fails.

Diagnostics available


Failure triage

Draft diagnoses alongside the existing disposition process.

Proof

Faster correct decisions with traceable evidence.

What AI contributes, from requirement to release



AI assists the testing team. For code, fixtures, failure evidence and release handoffs, [open the banking engineering walkthrough](#) →