

A shared QA platform for Harbor

A payment-release story for a fictional fintech with 75 offshore QA staff.

75


offshore QA staff

One payment release.

A platform the next team can reuse.

The assumed QA organization


Five application squads share 75 offshore QA staff. Each person appears in one primary-role group.



The pilot draws eight QA staff from this population. Headcount alone does not establish a savings opportunity.

One timeout can become two debits

PAY-142 adds safe retries to a CAD 100 transfer. The customer retries after a timeout.




Illustrative outcomes for initial balances of CAD 1,000 and CAD 0. A journal contains one debit and one credit. No fees, FX or other activity.

The product requirement is one transfer, one balanced journal and one customer-visible result. This is an authored scenario.

Where the release work goes

Assumed hands-on effort for one bounded release test pack. Eight stages add to 300 hours.




300 person-hours for this pack. Each activity belongs to one stage. Baseline already includes its ordinary review and rework.

The pack spans 10 business days. Waiting time and human effort use different units and must be measured separately.

The platform investment

Common context, fixtures and evidence serve every squad. Application teams retain business ownership.



Shared foundation: synthetic data · dependency stubs · artifact IDs · ownership · model configuration

Existing test frameworks and pipelines execute approved work. AI supports preparation and interpretation.

AI assistance across the release

WORK	PROPOSED ASSISTANCE
Prepare	Clarify requirements, draft scenarios, prepare data and diagnose environment readiness.
Deliver	Draft framework tests, recommend execution scope and summarize failures.
Decide	Review money-movement evidence and resolve outstanding defects before release.

The eight-stage walkthrough contains inputs, outputs, technology examples and a named human check for each stage.

Delivery maturity changes the first move

The same tool creates different work in a shared manual environment and a repeatable delivery system.

01

Shared & manual

Begin with requirements, test design and evidence drafts. Establish stable fixtures and repeatable runs before expanding generation.

02

Mixed maturity

Pilot payment APIs and one web journey. Add synthetic fixtures, provider stubs and evidence links alongside AI assistance.

03

Repeatable delivery

Extend generation and triage across more application teams. Evaluate test selection in shadow mode before changing required coverage.

Cloud hosting alone does not establish deployment, data-reset or test maturity. Profiles are illustrative assumptions.

What 78 hours of gross capacity means

Mixed-maturity target: 300 hours becomes 222, including 40 hours of review.



Purple segment = review. Total assisted effort = 222h.

78h

gross capacity
300 – 222

66h

after operation
78 – 12

33h

usable capacity
 $66 \times 50\%$

15 packs

to recover 480h setup
rounded up, steady assumptions

Illustrative capacity arithmetic. Tool, cloud and vendor charges are excluded. This is not cash ROI or a measured AI effect.

Ownership within the offshore model

OWNER	ACCOUNTABILITY
Shared enablement	Own context adapters, templates, fixtures, runner integration and coaching.
Application QA teams	Own business scenarios, expected results, execution and defect disposition.
Product and release owners	Resolve requirement questions and accept the release evidence.
Delivery management	Provide repository access, overlap hours and an explicit backlog for freed capacity.

Reskill domain testers through paired reviews. Agree supplier incentives and reuse ownership before scaling.

A proposed twelve-week engagement

A measured pilot creates a reusable delivery pattern and a decision about the next applications.

Weeks 1–2

Workflow assessment

Observe a release pack. Establish effort, waits and application readiness.

Output: baseline + pilot scope

Weeks 3–8

Payment pilot

Connect existing tools. Run reviewed tests on comparable packs and measure the full workflow.

Output: reusable pattern + evidence

Weeks 9–12

Second-team trial

Transfer ownership. Verify reuse and agree which applications qualify next.

Output: rollout decision + backlog

A planning sequence, not a guaranteed schedule. Environment readiness or inconclusive results can extend a phase.

The evidence needed to expand

MEASURE	EXPANSION EVIDENCE
Quality	All mandatory payment assertions pass; no unexplained skips or weakened checks.
Effort	Include preparation, review, correction, retesting and platform operation.
Delivery	Report elapsed time, waiting reasons and failed-run recovery separately.
Reuse	A second squad runs the pattern and maintains it without the pilot authors.

A faster draft is insufficient. Compare tasks of similar scope and retain conventional work as a reference.

The senior-director conversation

1. “Across the 75 QA staff, where do people spend time and where are they waiting?”

2. “Which payment journey has stable requirements, observable outcomes and an owner who can join a pilot?”

3. “Could one platform team remove duplicated work across the squads?”

4. “If we return capacity, which backlog or release commitment will use it?”

Proposed next step: a workflow assessment with the QA, development and platform leads. Agree a bounded pilot after reviewing the baseline.