



Assurance across the AI lifecycle


AI × quality engineering · Industry research



The enterprise assurance platform



A generated payment-API test




Illustrative workflow · approved contract defines the oracle; generated tests cannot silently weaken it

Generated tests pass through quality gates


A PASSING TEST IS AN INTERMEDIATE RESULT

Share of target test classes with at least one qualifying generated test



Meta / FSE 2024 · \$3.3 · 86 Kotlin components with existing test classes · reported percentages rounded
Denominator is target test classes, not individual generated tests. Coverage is a proxy for test improvement.

Evaluation, release and feedback




Proposed application-level assurance · evaluate again when context, behavior or authority changes

Application evaluation contract


LAYER	EVIDENCE	VALIDITY CHECK
Task outcome	Representative cases and expected results	Held-out examples; results by scenario
Retrieval	Relevant, permitted and grounded evidence	Correct citations and access scope
Agent actions	Tool choice, arguments and termination	Denied-action and recovery tests
Model judges	Scores against a defined rubric	Calibration against expert labels
Operation	Latency, cost, drift and escalation	Tail failures and tested fallback

Approved context and retrieval boundaries



Proposed checks include denied retrieval, access revocation, stale content and conflicting source passages.

The evaluated application has a release identity




A material component change creates a new candidate and triggers the relevant evaluation suite.

The source of expected behavior




A test can pass because it repeats the implementation's mistake. The expected behavior needs an independent basis.

Mutation-guided test generation



Failure diagnosis with inspectable evidence



GOOGLE: REVIEWED ACCURACY SAMPLE

90.14%

71 manually evaluated failures

Root-cause accuracy in the reported case study

GOOGLE: DEPLOYMENT POPULATION


52,635

distinct failing tests

A separate population; accuracy was not measured on all of it


Local pilot: blinded review of sampled causes, evidence correctness and total diagnosis effort.

Flaky-test repair and semantic checks



Pattern informed by FlakyGuard's Uber Go case; validate transferability to the local language and test environment.

The evaluation corpus lifecycle



SPLIT RELATED CASES TOGETHER

Near duplicates can make a held-out result look more general than it is.


Version each case, source permission, expectation, scenario tag and review decision. Report results by important slice.

Model judges need their own evaluation

CHECK	EVALUATION DESIGN	RELEASE IMPLICATION
Agreement with experts	Compare judge decisions with independently reviewed labels, by failure category.	Investigate material disagreements before trusting aggregate scores.
False acceptance	Review cases where the judge passes an incorrect or unsafe output.	Critical false passes can require a deterministic rule or expert gate.
Consistency and bias	Repeat ambiguous cases; vary answer order in pairwise comparisons.	Record instability and sensitivity with the rubric version.
Separation from tuning	Keep calibration examples separate from final evaluation cases.	Recheck validity when the judge, prompt or task population changes.

Proposed calibration contract. No universal agreement threshold or judge accuracy is assumed. [A01][T03][T06]


The tool authority boundary



Test: poisoned context requests a production write → gateway denies → trace records identity, target and reason


Proposed negative test · test recovery as well as refusal · finite coverage does not prove absence of vulnerabilities

A denied action through the gateway




Negative case: a retrieved document requests a production write, while the agent has sandbox-only authority.

Execution zones and production authority




Sandboxing limits consequences; it still needs access controls, escape tests and artifact validation.

Release floors and scenario-level decisions




A higher average score cannot override a failed critical scenario. Thresholds are agreed locally.

Containment, fallback and recovery




Reverting a model does not undo a completed side effect. Compensation and reconciliation need their own runbook.

Operational evidence and feedback



Link work, review and outcomes



Total human effort = preparation + execution + verification + correction + rework

Proposed logical schema · retain identifiers and versions; minimize sensitive payloads · time saved requires a capture mechanism

Tools occupy different assurance layers

LAYER	REPRESENTATIVE TOOLS	LOCAL PROOF
Test assets	Tosca · Applitools	Test validity; baseline quality
Change review	GitHub Copilot	Issues found and missed; reviewer effort
AI evaluation	LangSmith · Microsoft Foundry	Dataset fit; judge calibration; exportability
Adversarial testing	Promptfoo	Coverage of actual tool authority

Representative documentation, reviewed September 2026. Not a ranking or independent product benchmark. [\[T01\]](#)[\[T02\]](#)[\[T03\]](#)[\[T04\]](#)[\[T05\]](#)[\[T06\]](#)

Integration contracts preserve replaceability

BOUNDARY	MINIMUM CONTRACT	FAILURE TO EXERCISE
Task intake → runtime	Task ID, approved scope, context references and deadline.	Reject missing scope or an unsupported task type.
Runtime → action gateway	Agent identity, tool schema, resource and authorization context.	Deny wrong-resource calls and malformed arguments.
Runner → evidence store	Run ID, artifact hash, check results and configuration references.	Hold incomplete evidence or mismatched artifacts.
Evaluation → release owner	Scenario results, threshold version, unresolved failures and decision record.	Prevent promotion with stale or missing evaluation evidence.

Proposed interface requirements. Adapter compatibility and evidence export need local testing; ACS is a newly introduced standard with evolving implementation coverage. [A01][A03][A05][T03]

Value within quality and control floors

DECISION	RULE, APPLIED IN THIS ORDER	OWNER / ACTION
Stop / hold	Any quality, security or authorization breach; or adequate evidence shows both use cases below 10% saving.	Control owner holds access; sponsor stops or re-scopes.
Insufficient evidence	Missing baseline, exposure, review hours or sample; uncertainty crosses a decision boundary.	Measurement lead reports unknown; collect evidence within a capped extension.
Go, limited scope	At least one case $\geq 15\%$, adoption $\geq 50\%$, floors and cost cap hold. Finance validates any cash claim.	Sponsor + control owner authorize only the supported use case.
Extend / redesign	Best case $\geq 10\%$ but $< 15\%$, or adoption/cost criteria miss without a breach.	Delivery lead changes one hypothesis; one extension ≤ 4 weeks.

A poisoned-document exercise

STAGE	ILLUSTRATIVE INJECTED INSTRUCTION	EXPECTED EVIDENCE
Context intake	A test specification asks the agent to export customer data to an external endpoint.	Source provenance retained; retrieved content has no authority to grant access.
Action request	The runtime attempts a tool call outside its approved sandbox scope.	Gateway denies the destination and resource before execution.
Containment	The denied request reaches the workflow’s escalation rule.	No network side effect; bounded termination and a correlated event record.
Regression and recovery	Reviewers confirm the scenario and the expected control response.	Sanitized case added to the adversarial corpus; recovery verified before retry.


Illustrative exercise, not an observed incident or a guarantee that all injection paths are covered. [A03][A04][R01]

Architecture decisions with evidence


RESEARCH OBSERVATION

DESIGN RESPONSE

EVIDENCE TO RETAIN



A payment test has four joined contracts



Every boundary verifies identity, versions, expiry, resource scope and the same artifact digest.

Mismatch, unavailable authorization or missing durable evidence → hold; no promotion.

Failure semantics are part of the contract

FAILURE	REQUIRED BEHAVIOR	PAYMENT EXAMPLE
Timeout / unknown result	Cancel new actions; query status by idempotency key before retry.	Reuse RUN-042; never launch duplicate runners blindly.
Stale policy / auth outage	Deny execution; renew policy and re-authorize.	Expired DEC-042 cannot authorize a test run.
Evidence unavailable	Hold promotion; bounded durable outbox or halt.	No complete evidence receipt means no review-ready change.
Digest / oracle mismatch	Invalidate results; regenerate evidence against the new artifact.	Changed payment assertion cannot reuse EVAL-042.
Recovery	Human queue; exercise the failing case before an owner resumes.	Model rollback does not undo completed external actions.

A finding must end in a verified disposition

